


Advice for Applying Machine Learning

Quiz, 5 questions

1
point

1.

You train a learning algorithm, and find that it has unacceptably high error on the test set. You plot the learning curve, and obtain the figure below. Is the algorithm suffering from high bias, high variance, or neither?



- ☒ High bias
- ☐ Neither
- ☐ High variance

1
point

2.

Suppose you have implemented regularized logistic regression

Advice for Applying Machine Learning

Quiz, 5 questions

recognition). However, when you test your hypothesis on a new set of images, you find that it makes unacceptably large errors with its predictions on the new images. However, your hypothesis performs **well** (has low error) on the training set. Which of the following are promising steps to take? Check all that apply.

- ☒ Try using a smaller set of features.
- ☐ Try decreasing the regularization parameter λ .
- ☐ Try evaluating the hypothesis on a cross validation set rather than the test set.
- ☒ Try increasing the regularization parameter λ .

1
point

3.

Suppose you have implemented regularized logistic regression to predict what items customers will purchase on a web shopping site. However, when you test your hypothesis on a new set of customers, you find that it makes unacceptably large errors in its predictions. Furthermore, the hypothesis performs **poorly** on the training set. Which of the following might be promising steps to take? Check all that apply.

- ☒ Try adding polynomial features.



Try decreasing the regularization parameter λ .



Try evaluating the hypothesis on a cross validation set rather than the test set.



Use fewer training examples.

1

point

4.

Which of the following statements are true? Check all that apply.



Suppose you are training a regularized linear regression model. The recommended way to choose what value of regularization parameter λ to use is to choose the value of λ which gives the lowest **test set** error.



Suppose you are training a regularized linear regression model. The recommended way to choose what value of regularization parameter λ to use is to choose the value of λ which gives the lowest **training set** error.



The performance of a learning algorithm on the training set will typically be better than its performance on the test set.



Suppose you are training a regularized linear regression model. The recommended way to choose what value of regularization parameter λ to use is to choose the value of λ which gives the lowest **cross validation** error.

1

point

5.

Which of the following statements are true? Check all that apply.



A model with more parameters is more prone to overfitting and typically has higher variance.



When debugging learning algorithms, it is useful to plot a learning curve to understand if there is a high bias or high variance problem.

Advice for Applying Machine Learning

Quiz, 5 questions



If a neural network has much lower training error than test error, then adding more layers will help bring the test error down because we can fit the test set better.



If a learning algorithm is suffering from high bias, only adding more training examples may **not** improve the test error significantly.



I understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

Learn more about Coursera's Honor Code

Guoneng Zhong

Submit Quiz

